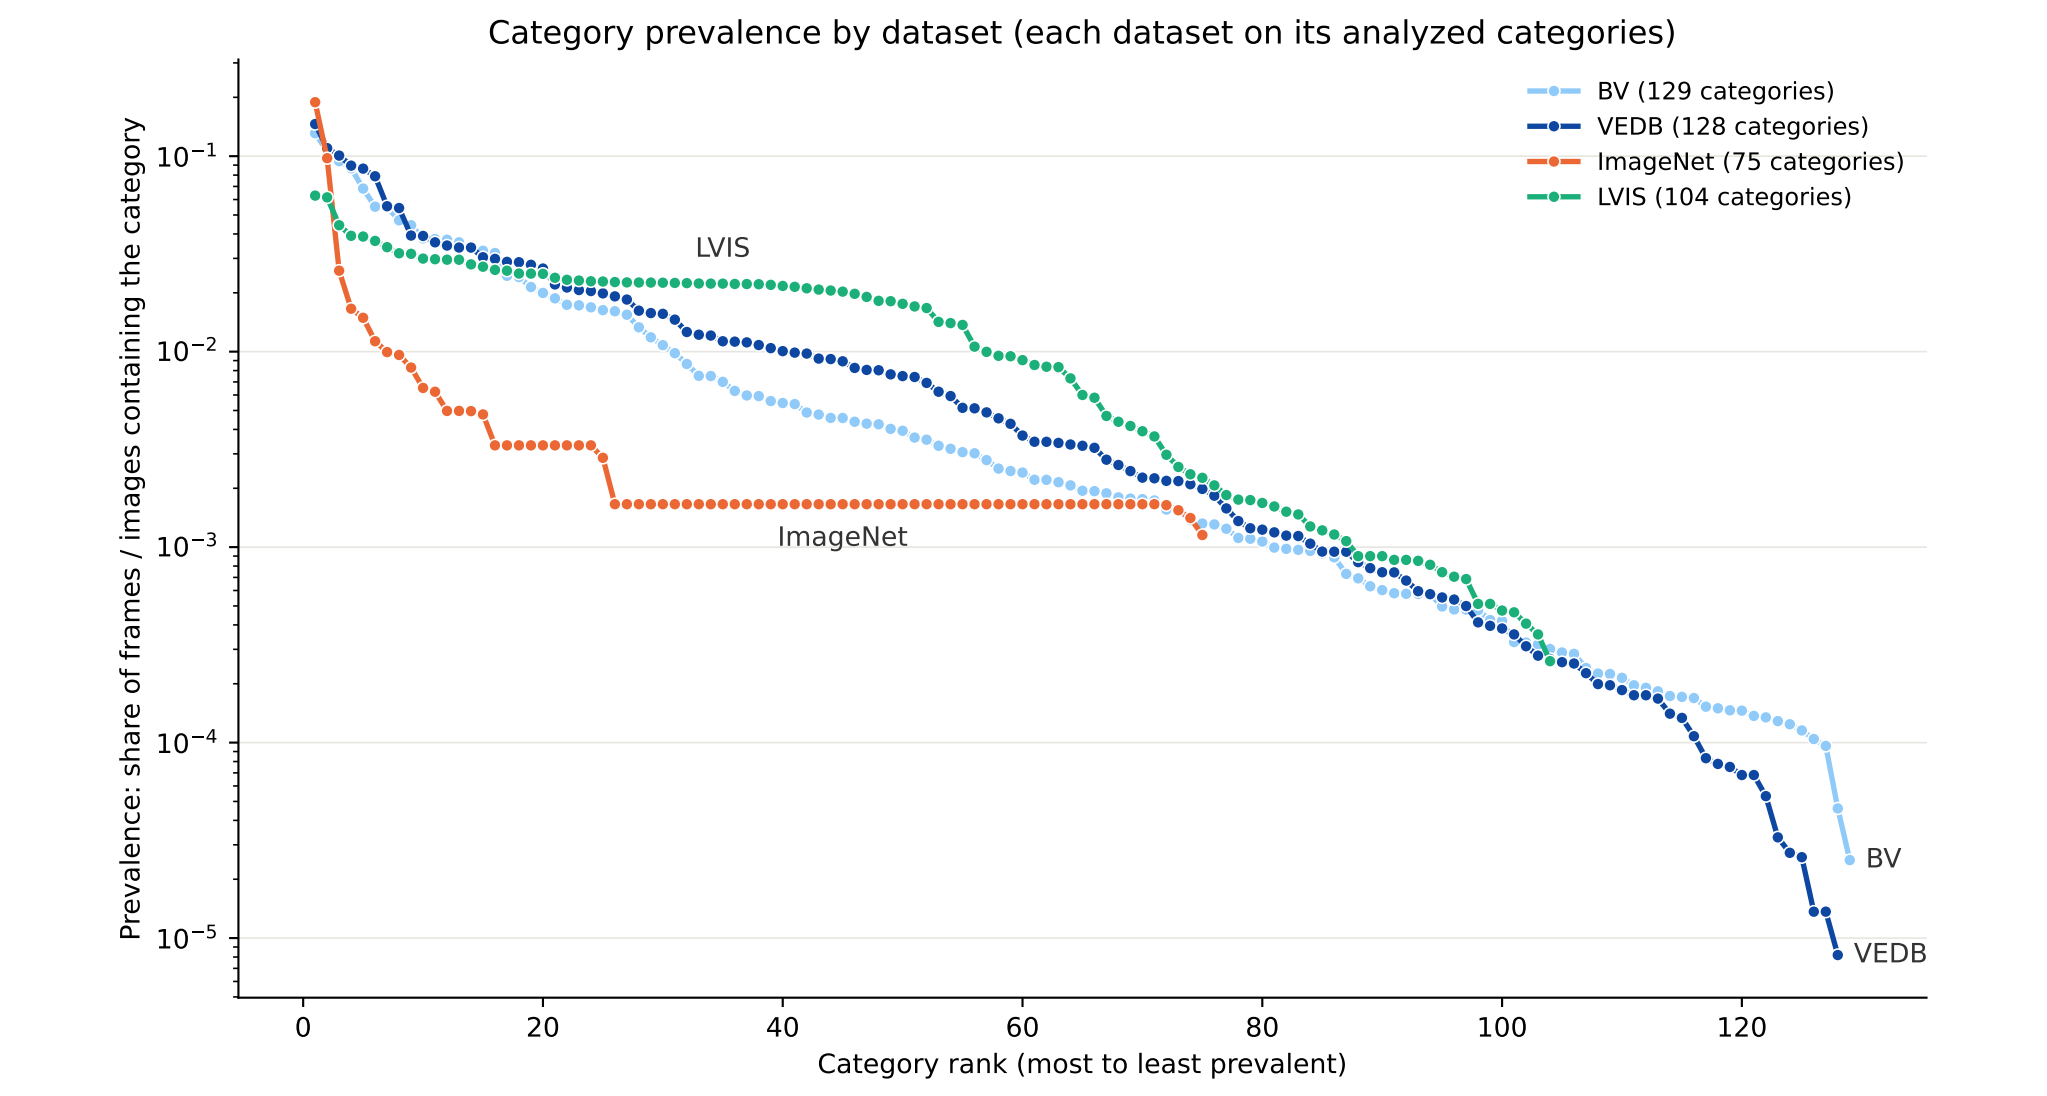

BV, VEDB: frames with a CLIP-passing detection. LVIS: images with a labeled box; LVIS caps the images labeled per category (≤2,364), so its curve reflects the labeling protocol. ImageNet: one label per image, ~1,300 images per class, so its curve mostly reflects how many classes each category pools.